

SCTP Sendbuffer Advertisement

A Thesis

submitted by

Arpan Kapoor B120555CS

Deepak Sirone J B120097CS

K Prasad Krishnan B120128CS

in partial fulfilment for the award of the degree of


Bachelor of Technology

in

Computer Science and Engineering

under the guidance of

Dr. Vinod Pathari



Department of Computer Science and Engineering
National Institute of Technology Calicut
NIT Campus P.O., Calicut
Kerala, India 673601

April 20, 2016

Abstract

TODO: UPDATE Advertising send buffer occupancy, i.e. the amount of backlogged data present in the sender's buffer has been proposed for TCP and experiments show improvements in application limited flows. This project proposes to advertise the same in SCTP.

Contents

1	Introduction	1
1.1	Problem Statement	1
1.2	Literature Survey	1
2	Design	3
2.1	Prerequisite terms	3
2.2	Proposed Modification	4
2.2.1	Send buffer value	4
2.3	Test bed design	5
3	Implementation	6
3.1	Linux Kernel Modification	6
3.2	Test bed Implementation Details	6
4	Results	8
4.1	Use Case Description	8
4.2	Test Results	8
4.3	Explanation of Test Results	8
5	Conclusions	9
5.1	Possible Quantitative Predictions	9

List of Figures

2.1	SCTP Packet Format	3
2.2	SCTP Chunk Format	4
2.3	Proposed Chunk for send buffer advertisement	4
2.4	Send buffer structure	5
2.5	Test bed topology	5
3.1	Test bed implementation	7

Chapter 1

Introduction

Stream Control Transport Protocol (SCTP) is a reliable transport protocol designed to transport Public Switched Telephone Network (PSTN) signaling messages over IP networks, but is capable of broader applications. Unlike TCP, SCTP offers sequenced delivery of user messages within multiple unidirectional logical channels called streams. Each SCTP endpoint is represented as a set of destination transport addresses, one of which is the primary address. If the primary address becomes unreachable SCTP chooses another destination transport address to route the messages thereafter. This provides network-level fault tolerance and is called multi-homing. It also employs a security cookie mechanism during association initialization to provide resistance to flooding and masquerade attacks.

TODO: Revise this paragraph. Advertising the amount of backlogged data present in the sender's buffer can help network operators evaluate the end-to-end performance of a connection in a better way than that with the existing passive measurements. This information can also be used to infer whether a connection is limited by the network or the application.

1.1 Problem Statement

TODO: Update. To propose a scheme to advertise send buffer occupancy information in SCTP, implement it in the Linux kernel and study the performance and security implications of the same.

1.2 Literature Survey

RFC 3286 [7] provides a high level introduction to the capabilities supported by SCTP, while RFC 4960 [8] describes the complete protocol. Agache and Raiciu

[1] propose a scheme to advertise send buffer occupancy in TCP. [4] was used to study the state machine employed in the Linux SCTP implementation. It was also used to understand the SCTP packet flow within the kernel. [5] provided with an overview of the traffic control and routing mechanisms in the Linux kernel, along with the userspace tools available for shaping and controlling the traffic.

Chapter 2

Design

2.1 Prerequisite terms

- **SCTP packet** is the unit of data that is passed on to the lower layer protocol. It is composed of a common header followed by one or more chunks.




Figure 2.1: SCTP Packet Format

- **SCTP Chunk** is a unit of information within an SCTP packet, containing either control information or user data. It consists of the following fields:

- **Chunk Type** identifies the type of information contained in the Chunk Value field. The chunk type values 16–62, 64–126, 129, 131, 133–190, 194–254 are currently unassigned[3].

The highest-order 2 bits of this 8-bit field specify the action that must be taken if the processing endpoint does not recognize the Chunk Type.

- 00 – Stop processing this SCTP packet and discard it.
- 01 – Stop processing this SCTP packet and discard it and report the unrecognized chunk in an ‘Unrecognized Chunk Type’.
- 10 – Skip this chunk and continue processing.

- 11 - Skip this chunk and continue processing, but report in an ERROR chunk using the ‘Unrecognized Chunk Type’ cause of error.
- **Chunk Flags** usage depends on the Chunk Type.
- **Chunk Length** represents the size of the chunk in bytes, including the Chunk Type, Chunk Flags, Chunk Length, and Chunk Value fields.
- **Chunk Value** contains the actual information to be transferred in the chunk.




Figure 2.2: SCTP Chunk Format

2.2 Proposed Modification

For advertising the send buffer occupancy, we propose to add a new Chunk Type, with a 32-bit Chunk Value field containing the amount of backlogged data in the send buffer.

To ensure that hosts running an unmodified SCTP stack skip this chunk without returning an ERROR chunk, the highest-order 2 bits of the Chunk Type of this chunk should be 10 (as explained in section 2.1). This limits the choice of the Chunk Type value between 128 and 190.

The total size of this chunk is 8 bytes, which is 0.53% of a typical 1500 byte packet.




Figure 2.3: Proposed Chunk for send buffer advertisement

2.2.1 Send buffer value

The kernel send buffer can be divided into 2 distinct parts:

- unacknowledged in-flight segments.
- segments waiting to be sent (backlog).

We propose to advertise the number of bytes in this backlog in the Chunk Value field of the proposed chunk.



Figure 2.4: Send buffer structure

2.3 Test bed design

A dumbbell shaped network topology was created with 2 routers in the center, and multiple hosts connected to each of the routers.




Figure 2.5: Test bed topology

Chapter 3

Implementation

3.1 Linux Kernel Modification

- A patch implementing the SCTP send buffer advertisement was created for Linux kernel `v4.6-rc4`.
- The send buffer advertisement chunk type value was set to 150.
- To modify the frequency at which send buffer advertisement chunks are sent, a `sysctl` interface was created. The default value was set to 5 seconds.
- A kernel timer was added corresponding to each SCTP association (within the `struct sctp_association`).
- A state table was created for this chunk, specifying the states in which the send buffer advertisement chunk should be generated and sent.

3.2 Test bed Implementation Details

Two Raspberry Pis were configured as routers and laptops were used as end hosts to achieve the dumbbell shaped topology.




Figure 3.1: Test bed implementation

Chapter 4

Results

4.1 Use Case Description

4.2 Test Results

4.3 Explanation of Test Results

Chapter 5

Conclusions

5.1 Possible Quantitative Predictions

Bibliography

- [1] A. Agache and C. Raiciu. *TCP Sendbuffer Advertising*. Internet-Draft draft-agache-tcpm-sndbufadv-00.txt. IETF Secretariat, July 2015.
- [2] Alexandru Agache and Costin Raiciu. “Oh Flow, Are Thou Happy? TCP Sendbuffer Advertising for Make Benefit of Clouds and Tenants”. In: *7th USENIX Workshop on Hot Topics in Cloud Computing (HotCloud 15)*. Santa Clara, CA: USENIX Association, July 2015. URL: <https://www.usenix.org/conference/hotcloud15/workshop-program/presentation/agache>.
- [3] Internet Assigned Numbers Authority. *Stream Control Transmission Protocol (SCTP) Parameters*. 2015. URL: <https://www.iana.org/assignments/sctp-parameters/sctp-parameters.xhtml#sctp-parameters-1> (visited on 04/20/2016).
- [4] Karthik Budigere. “Linux Implementation Study of Stream Control Transmission Protocol”. In: *Proceedings of Seminar on Network Protocols in Operating Systems*, p. 22.
- [5] Bert Hubert. *Linux Advanced Routing & Traffic Control HOWTO*. 2012. URL: <http://lartc.org/lartc.html> (visited on 04/20/2016).
- [6] M. Tim Jones. *Better networking with SCTP*. Feb. 28, 2006. URL: <http://www.ibm.com/developerworks/library/1-sctp/> (visited on 04/20/2016).
- [7] L. Ong and J. Yoakum. *An Introduction to the Stream Control Transmission Protocol (SCTP)*. RFC 3286. RFC Editor, May 2002, pp. 1–10. URL: <http://www.rfc-editor.org/rfc/rfc3286.txt>.
- [8] R. Stewart. *Stream Control Transmission Protocol*. RFC 4960. RFC Editor, Sept. 2007, pp. 1–152. URL: <http://www.rfc-editor.org/rfc/rfc4960.txt>.
- [9] R. Stewart et al. *Sockets API Extensions for the Stream Control Transmission Protocol (SCTP)*. RFC 6458. RFC Editor, Dec. 2011, pp. 1–115. URL: <http://www.rfc-editor.org/rfc/rfc4960.txt>.